

GEORGE LEWIS

220 East 18th Street, New York, NY

ghl227@gmail.com

(203) 918-4859

Summary: Data scientist and software engineer experienced in designing and deploying statistical models, building efficient data pipelines, and creating tools to improve developer productivity.

EXPERIENCE

LendUp

2013 - Present

Head of Risk and Analytics

Principal Data Scientist

- Founded the Risk and Analytics Team, later the Data Science Team. Designed and implemented the risk and underwriting program, growing company's loan portfolio by a factor of 200 while decreasing loss rates and achieving unit profitability. Led the team from seed round funding through series B.
- Designed, trained, deployed, and monitored statistical models using logistic regression, generalized additive models, decision trees, and random forests to evaluate the credit risk of all applicants and to identify and mitigate fraud.
- Built and leveraged analytics capabilities to ensure model accuracy, to track business metrics, to reveal trends in the loan portfolio, and to uncover opportunities for the risk program.
- Grew the Data Science Team to more than 10 members. Established best practices for building statistical models and performing analytics. Enforced high software standards and helped develop a culture of strong programming practices within the Data Science Team.
- Built production-quality software in Java, Python, and Clojure to implement the risk program, serve and score models statistical models, and to ingest and process data in real time.
- Recruited and trained the original Data Engineering Team. Collaborated with Data Engineering Team to build an ETL framework, jobs, and pipelines for supporting machine learning and analytics using Python, Scala, Airflow and Spark. Developed tooling to enable streamlined model deployment for data scientists.

ATLAS Experiment, CERN

2008 - 2013

Graduate Researcher on the Large Hadron Collider (LHC)

- National Science Foundation US LHC Graduate Student Support Award recipient.
- Built sophisticated models using custom likelihood functions and modern frequentist inference techniques to produce world accurate measurements of fundamental particle properties.
- Developed and maintained a statistical framework in C++ that implemented novel modeling techniques and was used extensively throughout a 3,000 person collaboration.
- Used parallel batch computing to analyze petabytes of data distributed worldwide across data centers.
- Wrote production C++ code to select interesting collisions of the 10 million created per second.
- Enforced coding standards, conventions, and best practices across a large C++ and Python code base.

EDUCATION

Ph.D. in Experimental High Energy Particle Physics

2013

New York University

B.A. in Physics and Mathematics

2007

Columbia University

SKILLS

Modeling: Linear models, decision trees, ensembles, boosting, neural networks, MCMC, hypothesis testing and confidence intervals, custom likelihoods and Bayesian modeling.

Languages: Python, Java, SQL, Scala, Clojure, C++. Some experience with R, C, Cython.

Tools: Scikit-Learn, Pandas, Scipy, Jupyter, Postgres, Redshift, BigQuery, Spark, Airflow, AWS, Weka, XGBoost, Matplotlib, Seaborn, Flask, Play, Ring, Tensorflow, Tidyverse, PyMC3.

For a sample of my work and code, see: <https://spontaneoussymmetry.com/work>